

Deep Learning : Multiclass Classification

Que1: Download a ResNet 50 trained on the ImageNet classification dataset

Steps Involved (Common for Part a and b):

1. Exploring Dataset: Dataset used: STL-10

The STL-10 dataset is an image recognition dataset with:

- 10 classes: airplane, bird, car, cat, deer, dog, horse, monkey, ship, truck.
- Train Dataset Size: 5000 Shape: (5000, 3, 96, 96)
- Test Dataset Size: 8000 Shape: (8000, 3, 96, 96)
- Count of Images belonging to each class: 500 each


```
{'airplane': 500, 'bird': 500, 'car': 500, 'cat': 500, 'deer': 500, 'dog': 500, 'horse': 500, 'monkey': 500, 'ship': 500, 'truck': 500}
```
- Images are 96x96 pixels, color.
- 500 training images (10 pre-defined folds), 800 test images per class.

2. Data Preprocessing: Transformed the data with following operations

- `Resize((255,255))`,
- `CenterCrop(224)`,
- `Convert to Tensor`
- `Normalize([0.5,0.5,0.5],[0.5,0.5,0.5])`

2. Created a Train and Test Data Loaders each of batch size 100.

3. Visualizing a batch of data



4. Moved data and labels to GPU if available.

a) Use the features extracted from the last fully-connected layer and train a multiclass SVM classifier on STL-10 dataset. Report the following

- i. Accuracy, Confusion Matrix on test data.
- ii. ROC curve

1. Link of Google Colab Code File:

<https://colab.research.google.com/drive/1ROVDFwSNdWJj7jjHxUWwM0fIXmsedHE-?usp=sharing>

2. Used a pre trained Resnet50 model for feature extraction by removing the last fully connected layer of the model

FC layer removed: Linear-174 [-1, 1000] 2,049,000

Last layer of Resnet50 model after removing FC layer:

AdaptiveAvgPool2d-173 [-1, 2048, 1, 1] 0

Total params: 23,508,032

3. Used the features extracted from the pretrained Resnet50 model as X_train to train a new model for multiclass classification using SVM classifier.

Shape of X_train: (5000, 2048) indicating a total of 5000 data points with each image being represented as 2048 feature columns.

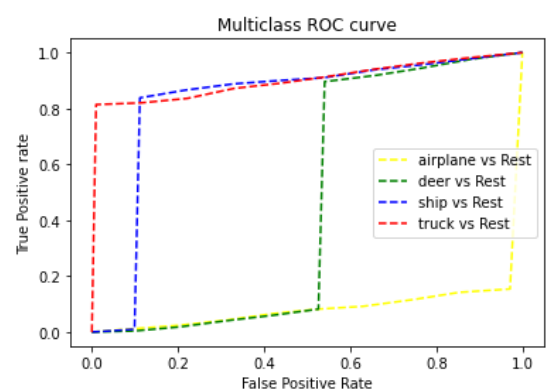
Length of Y_train: 5000, containing the class label for each of the images

SVM Classifier used to train the model:

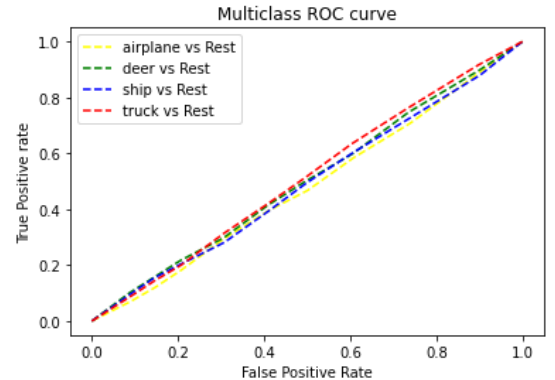
Kernel	Gamma	Regularizing factor (C)
'rbf'	0.001	200

4. Observation:

4.1. Train Data

Accuracy	Confusion Matrix	ROC Curve
82.66%	<pre>[[423 6 13 12 5 8 12 9 6 6] [8 433 9 9 5 9 8 7 9 3] [18 18 404 9 8 7 13 10 5 8] [17 9 8 427 8 6 8 6 9 2] [12 16 13 11 407 11 10 11 7 2] [9 12 13 18 10 408 12 8 5 5] [18 10 10 9 8 11 403 15 10 6] [27 11 6 12 5 8 15 407 2 7] [11 11 9 14 5 6 11 14 414 5] [9 10 11 14 11 9 18 8 3 407]]</pre>	 <p>Multiclass ROC curve</p>

4.2. Test Data

Accuracy	Confusion Matrix	ROC Curve
10.32%	<pre>[[88 146 98 94 58 81 88 51 40 56] [84 139 103 78 70 88 80 47 48 63] [78 117 106 91 68 77 82 60 38 83] [84 123 100 82 80 80 76 51 56 68] [82 121 117 79 70 93 65 49 44 80] [82 143 110 86 59 84 80 49 41 66] [70 142 101 87 70 84 79 56 42 69] [72 130 103 87 81 68 87 50 47 75] [96 124 98 87 80 90 62 41 47 75] [67 124 102 96 76 86 89 46 45 69]]</pre>	

b) Fine-tune the ResNet 50 model for the STL-10 dataset, and evaluate the classification performance on the test set before and after fine-tuning with respect to the following metrics

- Class wise Accuracy
- Report Confusion Matrix.

1. Link of Google Colab Code File:

<https://colab.research.google.com/drive/1jgnRkJkUu2K-8OeSVQN6zkdS6c4kW7r9?usp=sharing>

2. Observation:

2.1. Pretrained ResNet50 model (Before Fine Tuning)

- Accuracy on test data: 82.0%
- Class Wise accuracy:

Class	Accuracy
Airplane	84.25%
Bird	86.75%
Car	82.38%
Cat	77.75%
Deer	78.5%
Dog	71.88%
Horse	84.0%
Monkey	81.5%
Ship	90.75%
Truck	84.5%

2.2. Fine Tuned Resnet50 model:

Added following layers to the pretrained Resnet50 model and changed the output class from 1000 to 10 as the total classes present in the STL10 dataset is 10.

```
AdaptiveAvgPool2d-173      [-1, 2048, 1, 1]      0
  BatchNorm1d-174          [-1, 2048]            4,096
    Dropout-175            [-1, 2048]            0
      Linear-176            [-1, 2048]          4,196,352
        ReLU-177           [-1, 2048]            0
          BatchNorm1d-178   [-1, 2048]            4,096
            Dropout-179    [-1, 2048]            0
              Linear-180    [-1, 10]           20,490
=====
Total params: 27,733,066
Trainable params: 27,733,066
Non-trainable params: 0
-----
```

2.2.1. Trained the Fine tuned ResNet50 model:

No. of Epochs	Loss Function	Optimizer	Training Accuracy
10	CrossEntropyLoss	Adam	98.36%

```
Epoch: 0 Loss :0.2538050718545914 Accuracy :92.34%
Epoch: 1 Loss :0.12460703556537628 Accuracy :95.74000000000001%
Epoch: 2 Loss :0.09275018665790558 Accuracy :97.06%
Epoch: 3 Loss :0.06733751568389125 Accuracy :97.56%
Epoch: 4 Loss :0.048290332701802254 Accuracy :98.34%
Epoch: 5 Loss :0.0700740644901991 Accuracy :97.46000000000001%
Epoch: 6 Loss :0.060362133127450944 Accuracy :98.14%
Epoch: 7 Loss :0.05540539839863777 Accuracy :98.02%
Epoch: 8 Loss :0.048702784439921376 Accuracy :98.42%
Epoch: 9 Loss :0.04847560664564371 Accuracy :98.36%
Finished Training
```

2.2.2. Evaluation on Test Data

- Accuracy on Test data: 95.0%
- Class wise Accuracy:

Class	Accuracy
Airplane	97.5%
Bird	97.38%
Car	97.38%
Cat	95.88%

Deer	92.62%
Dog	91.75%
Horse	95.25%
Monkey	97.75%
Ship	95.62%
Truck	97.75%